

University of Toronto

Neural Networks for Machine Learning



Geoffrey Hinton - Summer 2015

Contributors: Max Smith

Latest revision: July 10, 2015

Contents

1	Introduction	1
2	The Perceptron learning procedure	1
3	The backpropagation learning procedure	1
4	Learning feature vectors for words	1
5	Object recognition with neural nets	1
5.1	Why object recognition is difficult	1
5.2	Achieving viewpoint invariance	1
	The invariant feature approach	1
	The judicious normalization approach	2
	The brute force normalization approach	2
6	Optimization: How to make the learning go faster	2
7	Recurrent neural networks	2
8	More recurrent neural networks	2
9	Ways to make neural networks generalize better	2

10 Combining multiple neural networks to improve generalization	2
11 Hopfield nets and Boltzmann machines	2
12 Restricted Boltzmann machines (RBMs)	2
13 Stacking RBMs to make Deep Belief Nets	2
14 Deep neural nets with generative pre-training	2
15 Modeling hierarchical structure with neural nets	2
16 Recent applications of deep neural nets	2

Abstract

Todo

1 Introduction

2 The Perceptron learning procedure

3 The backpropagation learning procedure

4 Learning feature vectors for words

5 Object recognition with neural nets

5.1 Why object recognition is difficult

- Things that make it hard to recognize objects:
 - **Segmentation:** Real scenes are cluttered with other objects
 - It's hard to tell which pieces go together as parts of the same object
 - Parts of an object can be hidden behind other objects
 - **Lighting:** The intensities of the pixels are determined as much by the lighting as by the objects
 - **Deformation:** Objects can deform in a variety of non-affine ways:
 - eg. a hand-written 2 can have a large loop or just a cusp
 - **Affordances:** Object classes are often defined by how they are used
 - **Viewpoint:** Changes in viewpoint cause changes in images that standard learning methods cannot cope with

5.2 Achieving viewpoint invariance

- Several approaches:
 - Use redundant invariant features
 - Put a box around the object and use normalized pixels
 - Use replicated features with pooling. This is called “convolution neural nets”
 - Use a hierarchy of parts that have explicit poses relative to the camera

The invariant feature approach

- Extract a large, redundant set of features that are invariant under transformations
- With enough invariant features, there is only one way to assemble them into an object
- But for recognition, we must avoid forming features from parts of different objects

The judicious normalization approach

- Put a box around the object and use it as a coordinate frame for a set of normalized pixels
 - This solves the dimension-hopping problem. If we choose the box correctly, the same part of an object always occurs on the same normalized pixels
 - The box can provide invariance to many degrees of freedom: translation, rotation, scale, shear, stretch, etc.
- Choosing the box is difficult because of: segmentation errors, occlusion, unusual orientations
- We need to recognize the shape to get the box right (chicken and egg dilemma)

The brute force normalization approach

- When training the recognizer, use well-segmented, upright images to fit the correct box
- At test time try all possible boxes in a range of position and scales
 - This approach is widely used for detecting upright things like faces and house numbers in unsegmented images
 - It is much more efficient if the recognizer can cope with some variation in position and scale so that we can use a coarse grid when trying all possible boxes

6 Optimization: How to make the learning go faster**7 Recurrent neural networks****8 More recurrent neural networks****9 Ways to make neural networks generalize better****10 Combining multiple neural networks to improve generalization****11 Hopfield nets and Boltzmann machines****12 Restricted Boltzmann machines (RBMs)****13 Stacking RBMs to make Deep Belief Nets****14 Deep neural nets with generative pre-training****15 Modeling hierarchical structure with neural nets****16 Recent applications of deep neural nets**